

Victoria University of Wellington  
School of Engineering and Computer Science

## SWEN430: Compiler Engineering (2018)

### Assignment 1 — Parsing and Interpreting Revised handout

Due: midnight, 12 August

#### Overview

This assignment will extend the parser and interpreter for the WHILE language which we have been studying in lectures.

Some of the features you are asked to add are described in the While Language Specification document (WLS), but not all of them. You are provided with around 100 `JUnit` tests cases to illustrate what is expected and help you determine whether the compiler is working correctly or not. If you are not sure what is expected, please ask for clarification.

**NOTE:** This assignment is not intended to handle union types, `null` or list/string concatenation (`++`). Tests which use string concatenation within `print` statements can easily be modified to use two `print` statements. Other tests including these features can be ignored.

#### The While Language Compiler

You can download an initial version of the compiler and the WHILE Language Specification from the SWEN430 homepage here:

[http://ecs.victoria.ac.nz/Courses/SWEN430\\_2018T2/Assignments/](http://ecs.victoria.ac.nz/Courses/SWEN430_2018T2/Assignments/)

To help get you started with the system, let's consider compiling the following file:

```
void main() {  
    print "Hello World!";  
}
```

If this file is stored in, say `test.while`, then you should be able to execute it using the built-in interpreter as follows:

```
% java test.while  
Hello World!
```

**NOTE:** This example changed — you are not provided with a `jar` file, so you need to compile the `While` compiler itself before you can run it.

As you can see, the output from the program is printed directly to the console. Furthermore, a `build.xml` file is provided to rebuild the `whilelang.jar` file after you have modified its source.

Passing the command-line switch `-verbose` will produce more verbose output which is useful for debugging. For example: `java -jar lib/whilelang.jar -verbose test.while`.

You can run the `JUnit` tests from within Eclipse by selecting “Run as JUnit Test” on the class `whilelang/testing/tests/RuntimeValidTests`.

**NOTE:** The given unit tests do not cover all aspects of the language, and you should supplement these with your own tests.

**NOTE:** There may be some mistakes in the test themselves — if you find any please let me know.

The changes you are required to make are described in the following sections. The changes are roughly in order of increasing complexity, but you don’t need to do them in the order listed.

## 1 Comments (10%)

Currently, the `WHILE` compiler does not support comments in the source code (see WLS, 2.2). You should find that at least two test cases (`SingleLineComment_Valid_1` and `MultiLineComment_Valid_1`) fail because of this. The objective here is to modify the compiler so that comments work as expected.

You might like to consider what the impact would be on the implementation if block comments (`/* ... */`) were allowed to be nested.

## 2 Real numbers (10%)

Several of the test files use a `real` type and real constants (e.g. `1.5`). Reals are used in the constant declaration example below, and are crucial to the examples of casts in part 6.

The objective of this part is to add support for real number. This will mean extending the scanner to recognise the keyword `real` and real constants; extending the parser and AST representation to handle these in declarations and expressions, respectively; and extending interpreter to store real values and apply arithmetic and comparison operators to them.

You should assume that a real constant consists of one or more digits followed by a decimal point (`“.”`), followed by one or more digits; i.e. there must be at least one digit before and after the decimal point. This means you can write `“1.0”`, `“0.001”`, `“121.5”`, etc; but you cannot write `“.1”`, `“5.”`, which are allowed in some languages.

You may optionally extend the type checker to check that real values are used correctly, but you are not required to do so.

## 3 Constant Declarations (10%)

The `WHILE` language is intended to support *constant declarations*, as follows:

```
const PI is 3.1415
int main() { print PI; }
```

When executed, this program should print “3.1415”, however at the moment you should find it does not, and that tests `Const_Valid_1 ... Const_Valid_4` fail. The objective here is to extend the `Interpreter` such that `const` declarations are appropriately implemented.

**HINT:** The simplest way to do this is to build a `Map<String,Expr>` from constant names to their bodies. Then, when an unknown `Expr.Variable` is encountered, you can check to see whether it’s actually referring to a `const` declaration (or not) — but you need to think about where that test is done.

An interesting extension here would be to allow the right hand side to be an expression containing only constants (including named constants, e.g. `2*PI`), rather than just constant values.

**NOTE:** You should assume that constant declarations can only occur at the “top level” of programs (as in the above example, and the test files mentioned above), not within methods.

## 4 Switch Statements (20%)

Currently, `switch` statements are not implemented in the compiler (see WSL, 5.3.3), and you should find that the tests `Switch_Valid_1 ... Switch_Valid_8` fail. The objective here is to implement `switch` statements appropriately. This will require modifying both the `Parser` and `Interpreter` classes. The syntax for `switch` statements is identical to that found in Java.

**HINT:** The `default` case can appear at any position in the `switch` statement, but is only taken if none of the other cases match. Furthermore, there may be at most one `default` case.

**HINT:** Java `switch` statements permit *unevaluated constant expressions* (see JLS§15.28). For example, one can write `case 1+1:` instead of `case 2:` and the Java compiler will evaluate this at compile time. **For simplicity, you may assume case conditions are either constants or named constants (i.e. from `const` declarations).**

**NOTE:** If a `break` statement is executed within a `switch` case, execution of the `switch` should be terminated and control passed to the statement following the `switch`. If execution reaches the end of a `switch` case without executing either a `break` or `return` statement, execution should continue with the next case in the `switch` statement. You may chose to ignore this “fall through” behaviour in your initial implementation.

**NOTE:** Java syntax allows a `default` case to occur anywhere in a `switch` statement. For this assignment, you may assume that if a `default` case is included, it must be the last case in the `switch` statement.

**NOTE:** You may optionally extend the type checker to check that `switch` statement satisfy the obvious rules (case values must have the same type as the `switch` expression, and cannot be repeated), but you are not required to do so.

## 5 String Assignments (20%)

Currently, the `WHILE` compiler does not support assignment to string elements. For example, the following code fails with a run time error:

```
void main() {
    string s = "Hello";
    s[0] = 'h';
}
```

```

    print s;
}

```

Similarly, the test `Char_Valid_3` fails because it assigns to a string element. The problem manifests itself in `Interpreter.execute(Stmt.Assign, HashMap<String, Object>)`, and is surprisingly tricky to resolve. Superficially, it appears as though the method simply doesn't consider the case of assignment to a string element (which it doesn't). Unfortunately, the problem is more involved because `Strings` in Java, unlike e.g. `ArrayLists`, are *immutable*. In contrast, `strings` in `WHILE` are *mutable*.

There are many ways you can go about resolving this issue. Perhaps the simplest is to move away from representing `WHILE strings` as Java `Strings`, but instead as e.g. `char[]` or similar. No matter how you address it, this will likely require widespread changes throughout the `Interpreter` to implement. Care should be taken to ensure that existing test cases still pass after you have fixed this problem.

## 6 Type Tests (15%)

Run time type tests are used to determine the type of a value at run time. The following illustrates a simple example:

```

bool isInt(int|null x) {
    if(x is int) { return true; }
    else { return false; }
}

```

Here, the `is` operator can be thought of as roughly equivalent to Java's `instanceof` operator. The `WHILE` language compiler does not currently implement this operator, and you should find that a number of test cases (`TypeEquals_Valid_1` ... `TypeEquals_Valid_16`) fail because of this. The objective is to implement the type test operator such that these tests will pass.

**NOTE:** As noted above, this assignment is not intended to handle unions or `null`, so the above example and test cases are not relevant. I will leave this in the assignment anyway, and you can implement a mechanism that could be used if we added unions later.

## 7 Explicit Casts (15%)

The `WHILE` language is intended to support *explicit coercions* (a.k.a *casts*), as follows:

```

int main() {
    int x = 1;
    real y = (real) x;
    print y;
}

```

Executing this method should print "1.0"; however, currently, you should find that it prints "1", which indicates that the coercion from an `int` to a `real` did not take place. Currently, the tests `Cast_Valid_1` ... `Cast_Valid_4` will fail.

The objective here is to implement **primitive casts**. Primitive casts can be used to convert an `int` into a `real`, as illustrated above, and can also be applied to structured types with `ints` as components. For example, we can cast from `[int]` to `[real]` by performing the primitive cast on all elements. Likewise, `{int x}` can be cast to `{real x}`. Note that such casts require an explicit coercion at run time.

## Submission

Your program code should be submitted electronically via the *online submission system*, linked from the course homepage. You must ensure your submission meets the following requirements (which are needed for the automatic marking script):

1. Your submission is packaged into a jar file, including the source code (see the export-to-jar tutorial linked from the course homepage).
2. The names of all classes, methods and packages remain unchanged.
3. You have removed any debugging code that produces output, or otherwise affects the computation.

You should also submit short document giving a summary of what you did for the assignment. This should have a section for each of the 7 parts of the assignment listed above, stating whether that part was complete, attempted but not completed, or not attempted. For those parts you attempted, you should provide a brief summary of what you did, including any major problems you encountered, anything you found difficult or interesting, any test cases you added (or modified), any problems that remain in the code you have submitted, and anything extra you did beyond the basic requirements of the assignment.

This document must have your name at the top, along with the name of the course and the assignment. It should be either a plain text or pdf file (latex output is ideal), and probably about one to two pages long.

**Note:** Failure to meet these requirements could result in you getting zero marks for the assignment.

## Submission

The assignment will be marked primarily on the correctness of your code, as determined by automated testing and additional tests that you provide. You will not be marked directly on your project summary or on coding style, but you will lose marks if you do not provide an adequate summary of what you did, if I find problems in your code that you have not told me about (in which case I will assume that you didn't know what them), or if I spot any especially bad coding.